

# Docker 101

# Plan

- Why Docker?
- Terminology
- Ecosystem
- Dockerfile
- CLI
- Tips & Tricks
- Docker-compose
- Tutorial

# Why docker?



Data Scientist



SuperComputer

# Why docker?



Data Scientist



SuperComputer

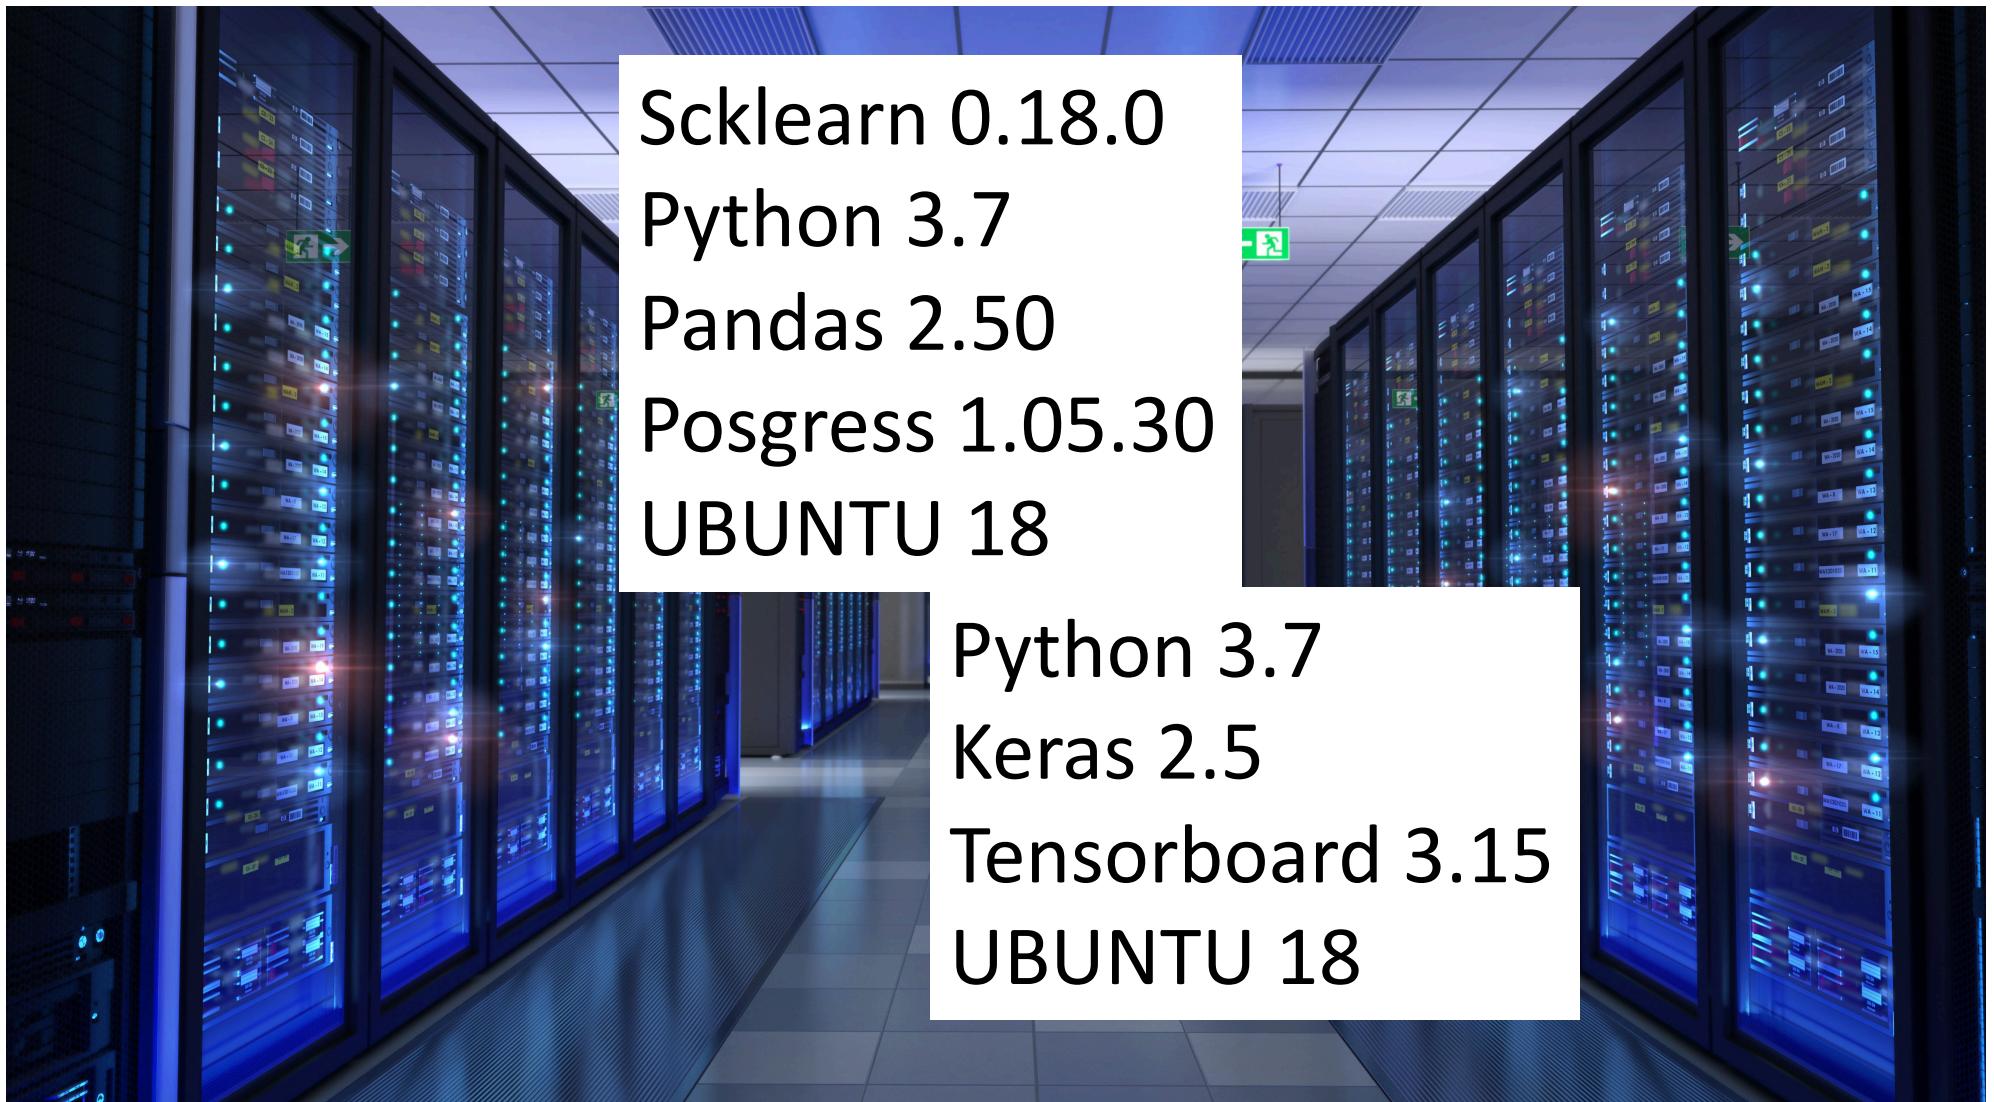
# Why docker?



Data Scientist



Data Scientist



SuperComputer

# Why docker?



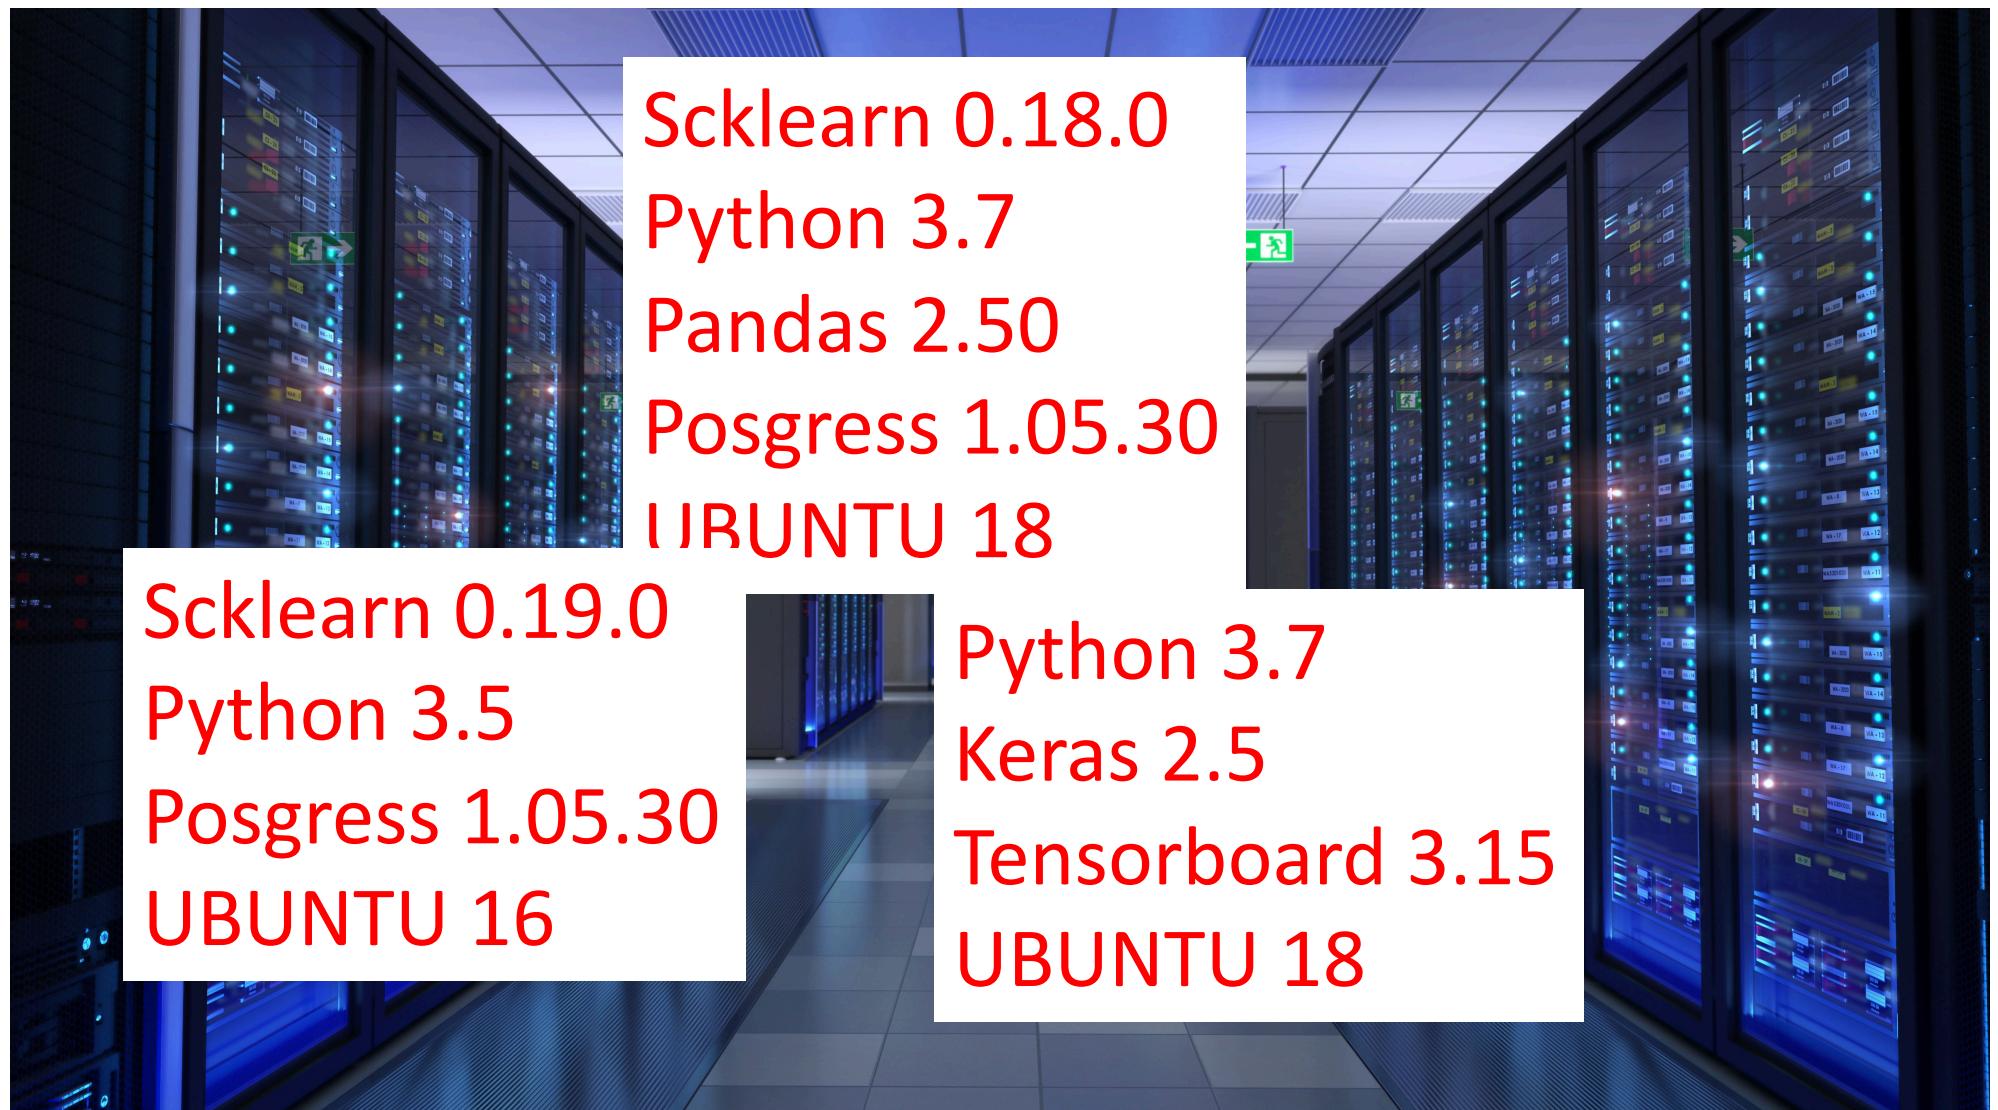
Data Scientist



Data Scientist



Data Scientist



SuperComputer

# Why docker?



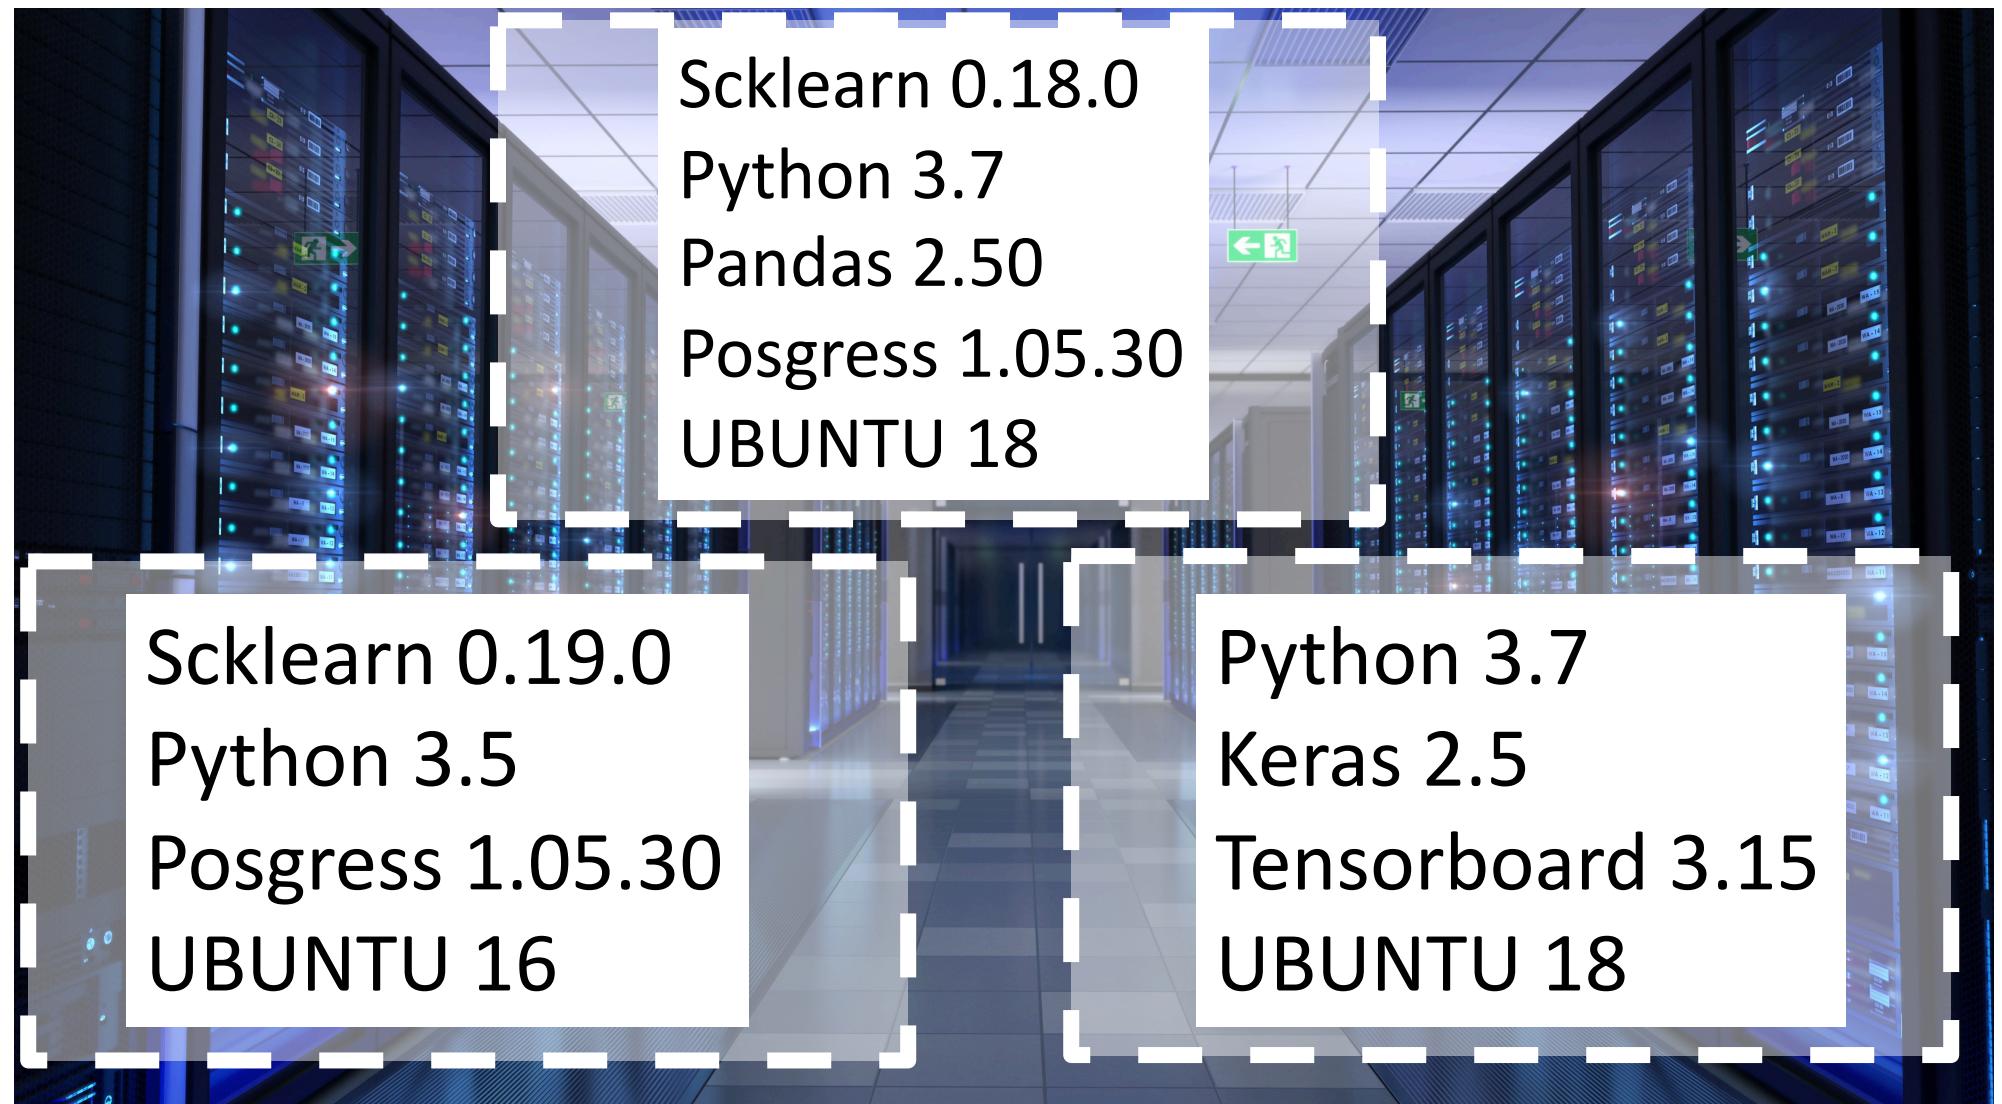
Data Scientist



Data Scientist

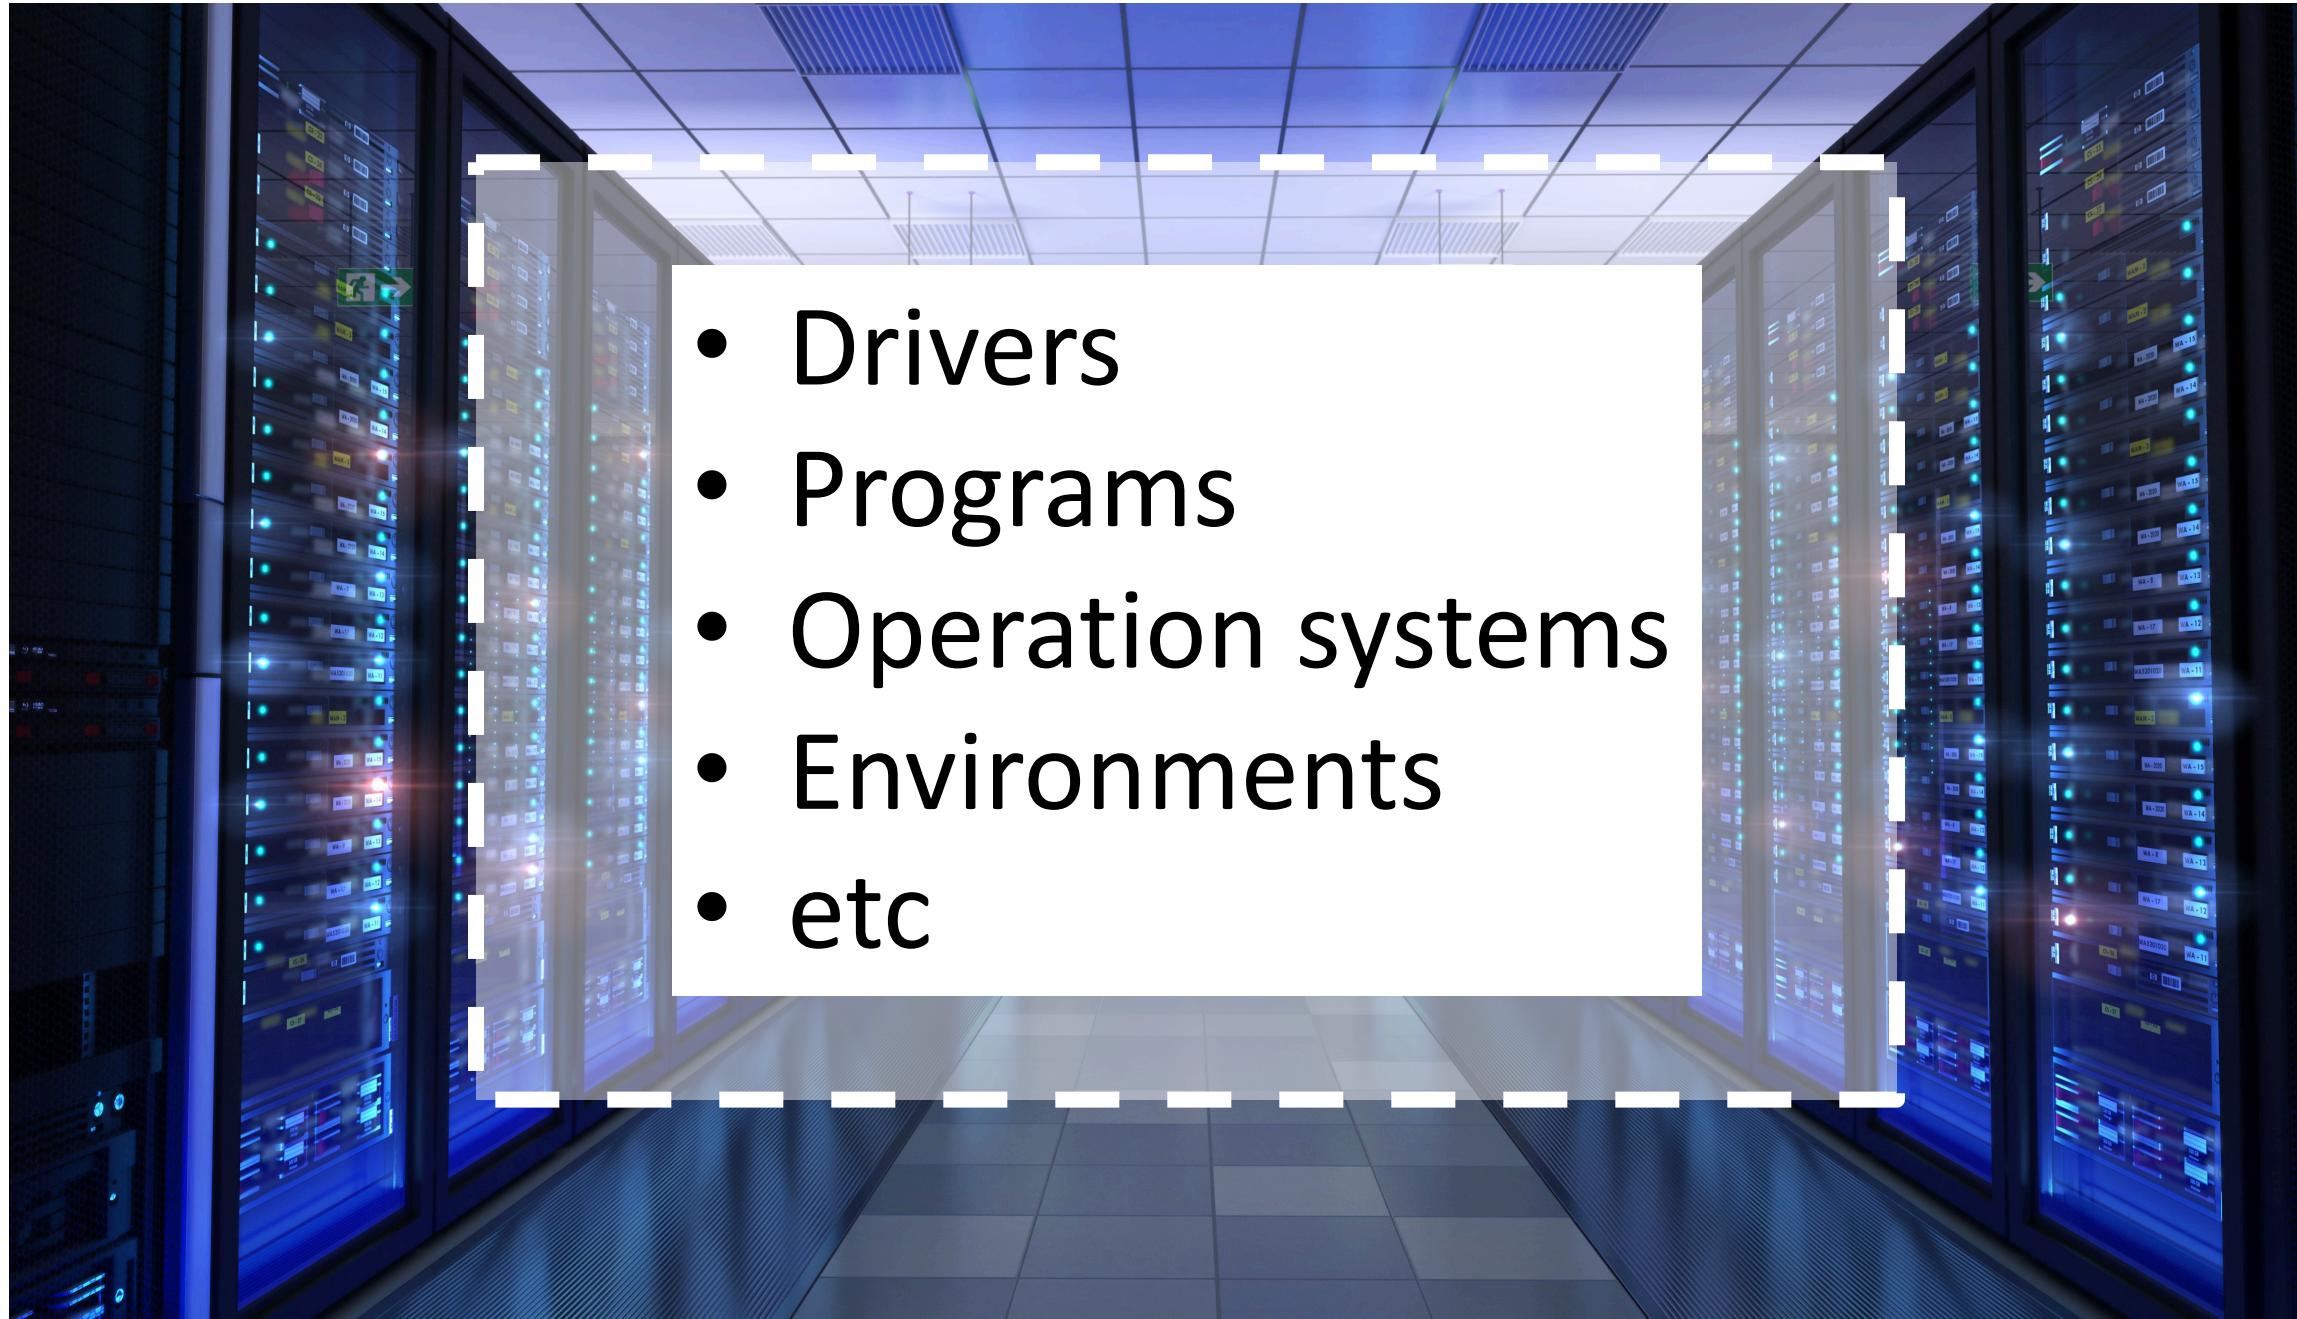


Data Scientist



SuperComputer

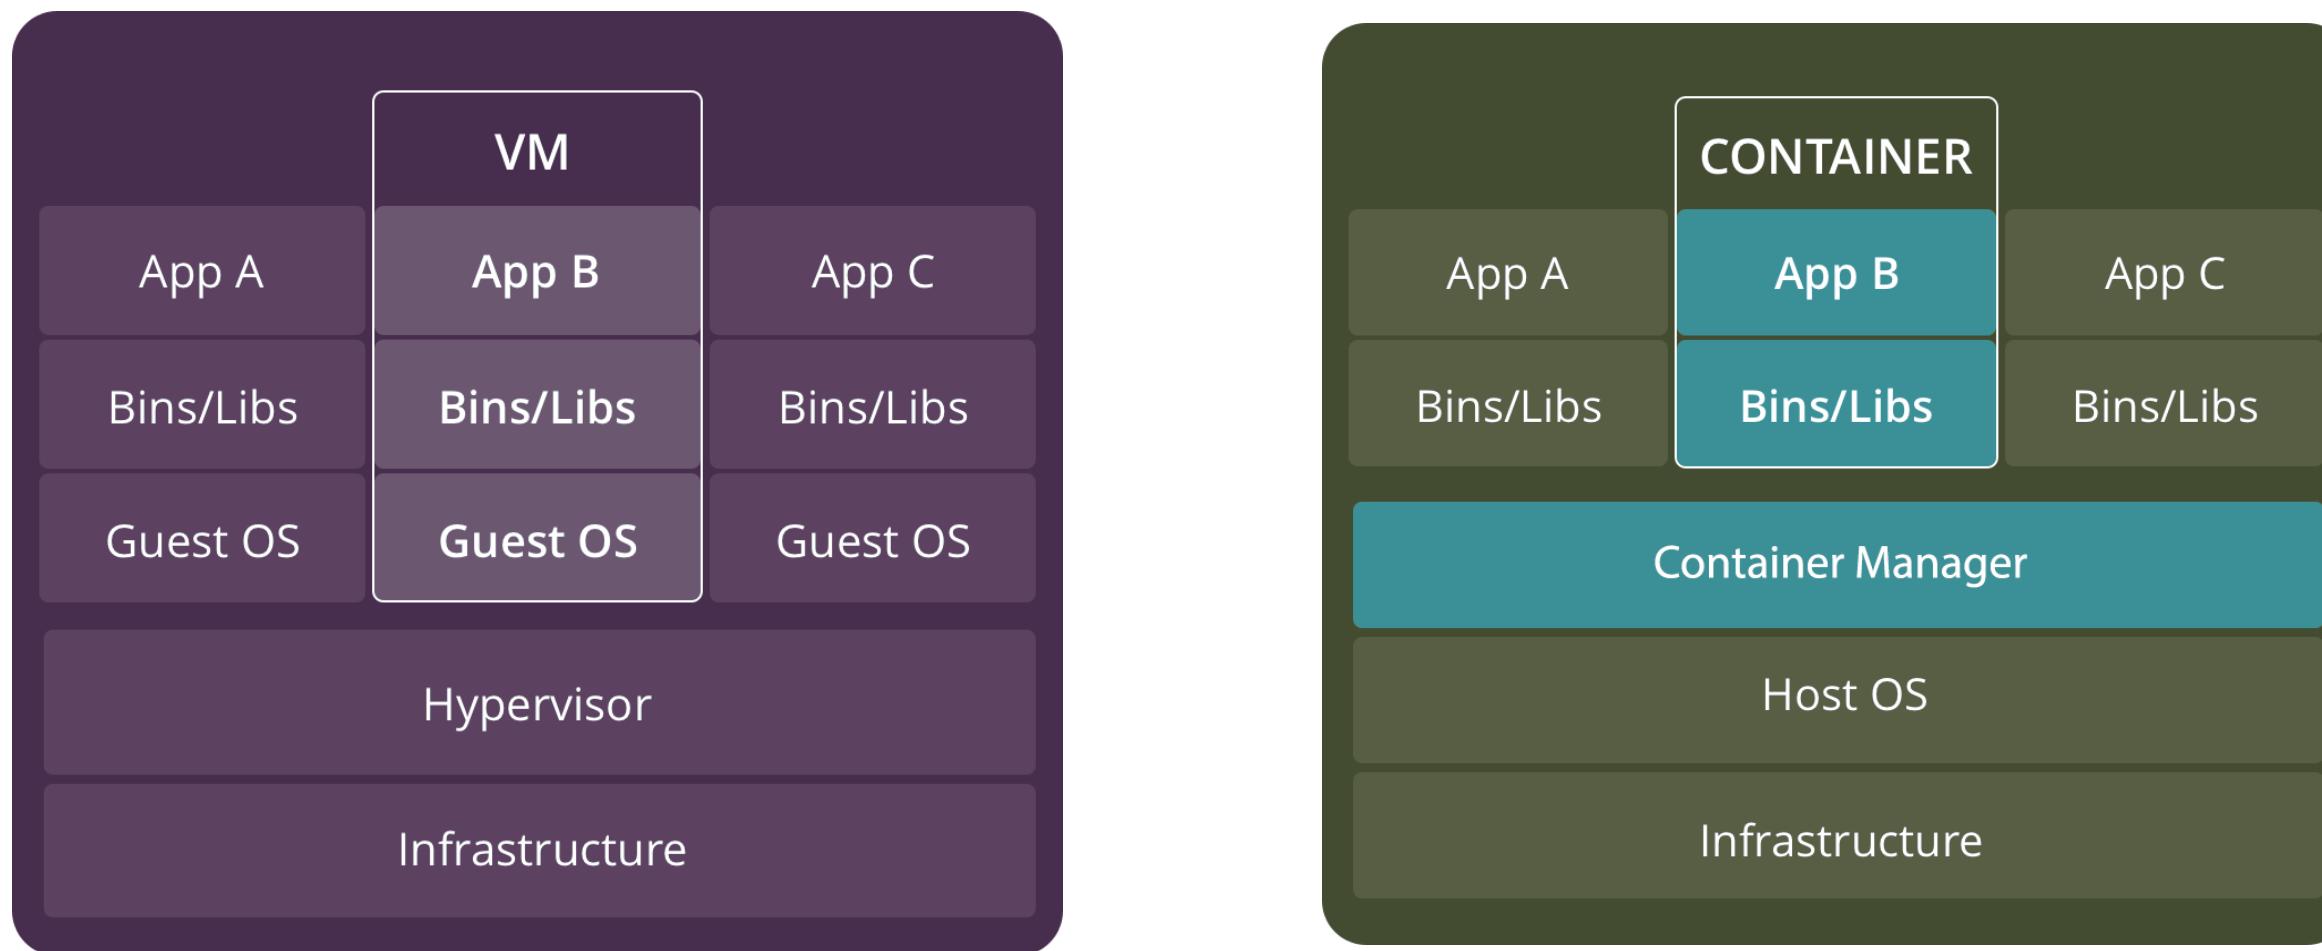
# Why docker?



- Drivers
- Programs
- Operation systems
- Environments
- etc

SuperComputer

# Why containers?



# Docker terminology

Dockerfile

Image

Container

---

---

# Docker terminology

Dockerfile

Image

Container

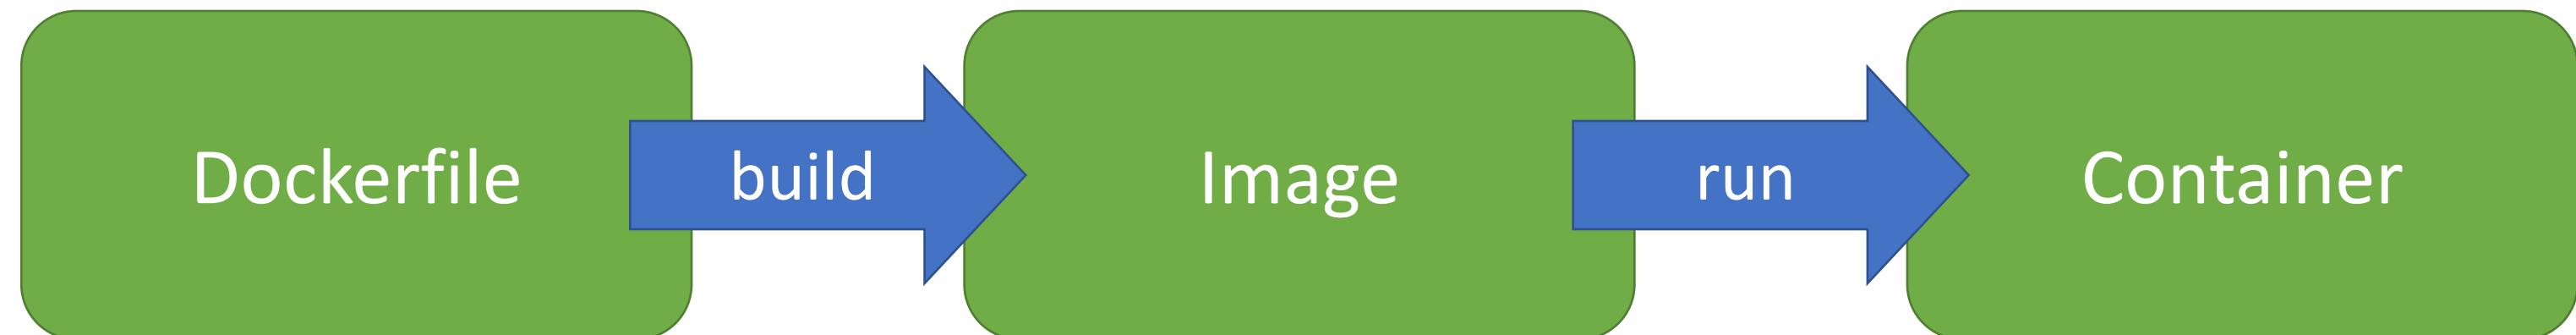
---

Formal class  
description

Class

Class instance  
(object)

# Docker terminology



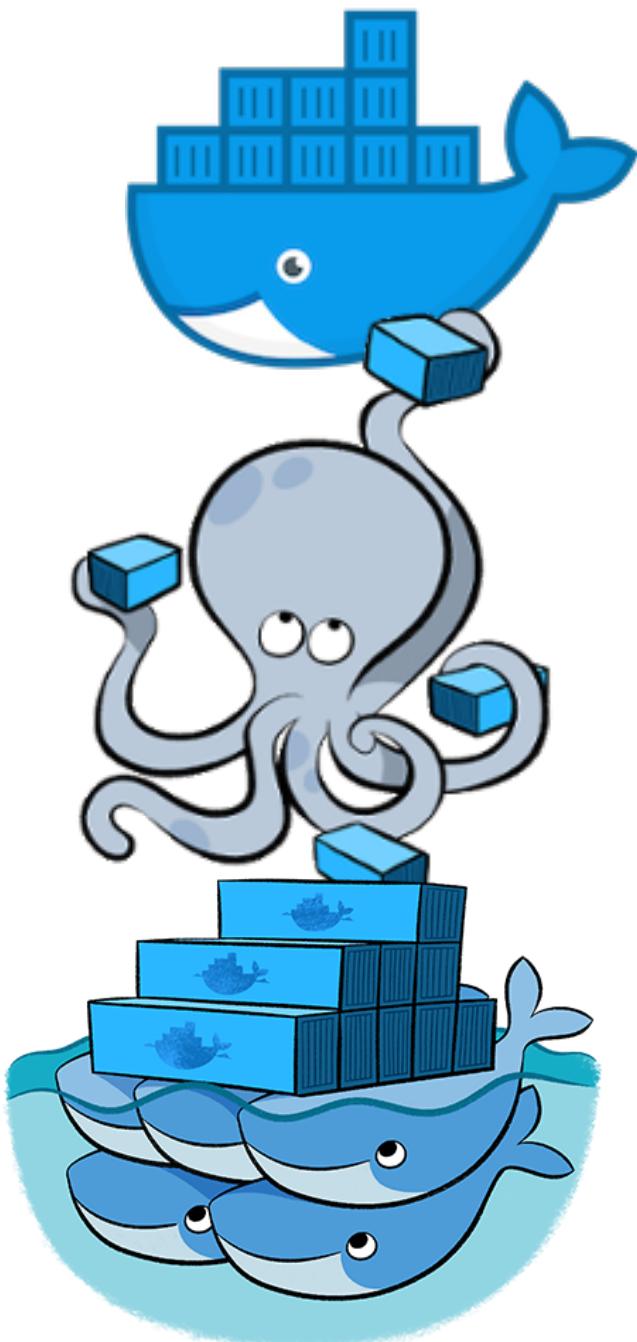
Formal class  
description

Class

Class instance  
(object)

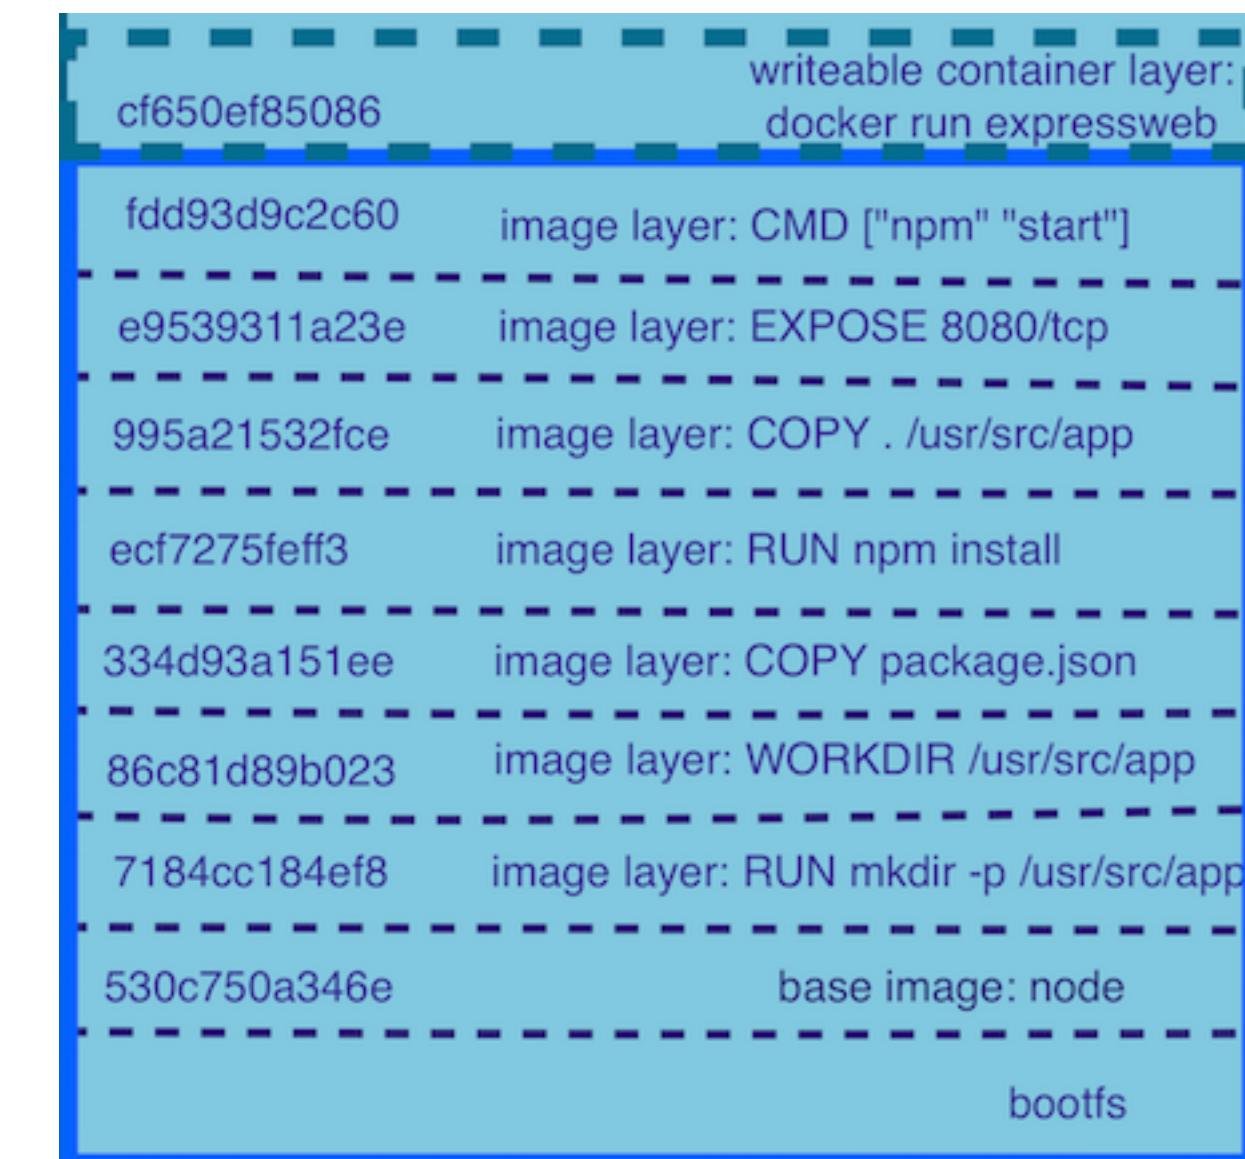
# What is docker?

- Docker Containers
- Docker Hub
- Docker Compose
- Docker Swarm
- Docker Cloud



# Dockerfile reference

- FROM
- RUN
- ENTRYPOINT / CMD
- COPY / ADD
- WORKDIR



# Docker cli basic

```
docker build -t alikhil/my-image:v1 -f Dockerfile .
```

```
docker run -p 8080:80 -d alikhil/my-image:v1
```

# Logs

```
docker logs [-f] [--tail=N] cntr_id_or_name
```

```
docker logs -f nginx
```

## Exec

```
docker exec [--it] cntr_name_or_id command
```

```
docker exec --it myapp bash
```

# Volumes

```
docker run -v host_dir:cntr_dir image-name
```

# Port forwarding

```
docker run -p <host-port>:<container-port> image-name
```

Example stateful postgres in docker:

```
docker run -d --rm  
  -v $(PWD)/data/pg:/var/lib/postgresql/data  
  -e POSTGRES_PASSWORD=1234  
  -e POSTGRES_USER=postgres  
  -p 5432:5432  
  --name pg-app postgres
```

# Real example

```
docker run --runtime=nvidia \
-e NVIDIA_VISIBLE_DEVICES=<Parameter> \
-e LOGIN=$(whoami) \
--cap-add SYS_ADMIN \
--cap-add DAC_READ_SEARCH \
-e LOGIN_UID=$(id -u) \
-e LOGIN_GID=$(id -g) \
-v $HOME/.ssh/authorized_keys:/mnt/authorized_keys:ro \
-v $HOME/work:/work/work \
--name $(whoami) \
-it \
--memory 64g \
-p 8888 \
-p 6006 \
-p 22 \
registry.skbkontur.ru/research/anaconda3-cuda10.0
```

# Trick #1

Speed up builds by:

- caching layers
- use .dockerignore

```
1  FROM node:8.12
2
3  RUN mkdir app
4  WORKDIR /app
5
6  ADD package.json
7
8  RUN npm install
9
10 ADD .
11
12 CMD ["node", 'index.js']
```

<https://blog.playmoweb.com/speed-up-your-builds-with-docker-cache-bfed14c051bf>

# Trick #2

Decrease image size by:

- using alpine images:
- combining several RUN layers into one

<https://hackernoon.com/tips-to-reduce-docker-image-sizes-876095da3b34>

# Trick #3

Save disk space in your host machine by:

- running containers with `--rm` argument
- or remove them later with:

`docker rm $(docker ps -q -f 'status=exited')`

- prune old volumes `docker volume prune`

- delete unused images - `docker rmi $(docker images -q -f "dangling=true")`

# Trick #4

Docker run can become very complicated. Like this:

```
docker run -d --rm  
  -v $(PWD)/data/pg:/var/lib/postgresql/data  
  -e POSTGRES_PASSWORD=1234  
  -e POSTGRES_USER=postgres  
  -p 5432:5432  
  --name pg-app postgres
```

Solution: use `docker-compose`

# docker-compose

```
version: '3'

services:
  redis:
    image: redis:5.0
    volumes:
      - $PWD/redis_data:/data
    ports:
      - 6379:6379

  postgres:
    image: postgres:9.6
    environment:
      - POSTGRES_USER=postgres
      - POSTGRES_PASSWORD=1234
      - POSTGRES_DB=postgres
    ports:
      - 5433:5432
    volumes:
      - $PWD/postgres_data:/var/lib/postgresql/data
```

# Resources

- <https://www.backblaze.com/blog/vm-vs-containers/>
- <https://github.com/alikhil/docker-101-tutorial>